

MathDNN Homework 11

Department of Computer Science and Engineering
2021-16988 Jaewan Park

Problem 1

(a) Since log is a concave function, using Jensen's inequality, we obtain the following.

$$\begin{aligned}
 \text{VLB}_{\theta, \phi}^{(K)}(x) &= \mathbb{E}_{Z_1, \dots, Z_K \sim q_\phi(z|x)} \left[\log \frac{1}{K} \sum_{k=1}^K \frac{p_\theta(x | Z_k) p_Z(Z_k)}{q_\phi(Z_k | x)} \right] \\
 &\leq \log \left(\mathbb{E}_{Z_1, \dots, Z_K \sim q_\phi(z|x)} \left[\frac{1}{K} \sum_{k=1}^K \frac{p_\theta(x | Z_k) p_Z(Z_k)}{q_\phi(Z_k | x)} \right] \right) \\
 &= \log \left(\frac{1}{K} \sum_{k=1}^K \mathbb{E}_{Z_k \sim q_\phi(z|x)} \left[\frac{p_\theta(x | Z_k) p_Z(Z_k)}{q_\phi(Z_k | x)} \right] \right) \\
 &= \log \left(\frac{1}{K} \sum_{k=1}^K p_\theta(x) \right) = \log p_\theta(x)
 \end{aligned}$$

(b) Using the given hint together with Jensen's inequality, we obtain the following.

$$\begin{aligned}
 \text{VLB}_{\theta, \phi}^{(K)}(x) &= \mathbb{E}_{Z_1, \dots, Z_K \sim q_\phi(z|x)} \left[\log \frac{1}{K} \sum_{k=1}^K \frac{p_\theta(x | Z_k) p_Z(Z_k)}{q_\phi(Z_k | x)} \right] \\
 &= \mathbb{E}_{Z_{i_1}, \dots, Z_{i_M} \sim q_\phi(z|x)} \left[\log \left(\mathbb{E}_{I=\{i_1, \dots, i_M\}} \left[\frac{1}{M} \sum_{m=1}^M \frac{p_\theta(x | Z_{i_m}) p_Z(Z_{i_m})}{q_\phi(Z_{i_m} | x)} \right] \right) \right] \\
 &\geq \mathbb{E}_{Z_{i_1}, \dots, Z_{i_M} \sim q_\phi(z|x)} \left[\mathbb{E}_{I=\{i_1, \dots, i_M\}} \left[\log \frac{1}{M} \sum_{m=1}^M \frac{p_\theta(x | Z_{i_m}) p_Z(Z_{i_m})}{q_\phi(Z_{i_m} | x)} \right] \right] \\
 &= \mathbb{E}_{I=\{i_1, \dots, i_M\}} \left[\mathbb{E}_{Z_{i_1}, \dots, Z_{i_M} \sim q_\phi(z|x)} \left[\log \frac{1}{M} \sum_{m=1}^M \frac{p_\theta(x | Z_{i_m}) p_Z(Z_{i_m})}{q_\phi(Z_{i_m} | x)} \right] \right] \\
 &= \mathbb{E}_{I=\{i_1, \dots, i_M\}} \left[\text{VLB}_{\theta, \phi}^{(M)}(x) \right] = \text{VLB}_{\theta, \phi}^{(M)}(x)
 \end{aligned}$$

(c) We should choose q_ϕ powerful enough so that $q_\phi(Z_k | x) = p_\theta(Z_k | x)$ for all $k = 1, \dots, K$.

$$\begin{aligned}
 \text{VLB}_{\theta, \phi}^{(K)}(x) &= \mathbb{E}_{Z_1, \dots, Z_K \sim q_\phi(z|x)} \left[\log \frac{1}{K} \sum_{k=1}^K \frac{p_\theta(x | Z_k) p_Z(Z_k)}{q_\phi(Z_k | x)} \right] \\
 &= \mathbb{E}_{Z_1, \dots, Z_K \sim q_\phi(z|x)} \left[\log \frac{1}{K} \sum_{k=1}^K \frac{p_\theta(x | Z_k) p_Z(Z_k)}{p_\theta(Z_k | x)} \right] \\
 &= \mathbb{E}_{Z_1, \dots, Z_K \sim q_\phi(z|x)} \left[\log \frac{1}{K} \sum_{k=1}^K p_\theta(x) \right] = p_\theta(x)
 \end{aligned}$$

Problem 2

(a) Since \log is a concave function, using Jensen's inequality, we obtain the following.

$$\begin{aligned}\log p_\theta(X_i) &= \log \left(\mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right] \right) \\ &\geq \mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\log \left(\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right) \right] = \text{VLB}_{\theta, \phi, \lambda}(X_i)\end{aligned}$$

(b) Gradients regarding θ and λ can be easily derived as the following.

$$\begin{aligned}\nabla_\theta \text{VLB}_{\theta, \phi, \lambda}(X_i) &= \nabla_\theta \mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\log \left(\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right) \right] = \mathbb{E}_{Z \sim q_\phi(z|X_i)} [\nabla_\theta \log p_\theta(X_i | Z)] \\ \nabla_\lambda \text{VLB}_{\theta, \phi, \lambda}(X_i) &= \nabla_\lambda \mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\log \left(\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right) \right] = \mathbb{E}_{Z \sim q_\phi(z|X_i)} [\nabla_\lambda \log r_\lambda(Z)]\end{aligned}$$

For gradients on ϕ , we can use the log-derivative trick.

$$\begin{aligned}\nabla_\phi \text{VLB}_{\theta, \phi, \lambda}(X_i) &= \nabla_\phi \mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\log \left(\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right) \right] = \nabla_\phi \int \log \left(\frac{p_\theta(X_i | z) r_\lambda(z)}{q_\phi(z | X_i)} \right) q_\phi(z | X_i) dz \\ &= \int \left(-\frac{\nabla_\phi q_\phi(z | X_i)}{q_\phi(z | X_i)} q_\phi(z | X_i) + \log \left(\frac{p_\theta(X_i | z) r_\lambda(z)}{q_\phi(z | X_i)} \right) \nabla_\phi q_\phi(z | X_i) \right) dz \\ &= \int \log \left(\frac{p_\theta(X_i | z) r_\lambda(z)}{q_\phi(z | X_i)} \right) \frac{\nabla_\phi q_\phi(z | X_i)}{q_\phi(z | X_i)} q_\phi(z | X_i) dz \\ &= \mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\log \left(\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right) \nabla_\phi \log q_\phi(Z | X_i) \right]\end{aligned}$$

(c) We can rewrite VLB as the following.

$$\begin{aligned}\text{VLB}_{\theta, \phi, \lambda}(X_i) &= \mathbb{E}_{Z \sim q_\phi(z|X_i)} \left[\log \left(\frac{p_\theta(X_i | Z) r_\lambda(Z)}{q_\phi(Z | X_i)} \right) \right] \\ &= \mathbb{E}_{Z \sim q_\phi(z|X_i)} [\log p_\theta(X_i | Z)] - D_{\text{KL}}(q_\phi(z | X_i) \parallel r_\lambda(z))\end{aligned}$$

Then the first term can be calculated as

$$\begin{aligned}\mathbb{E}_{Z \sim q_\phi(z|X_i)} [\log p_\theta(X_i | Z)] &= \mathbb{E}_{Z \sim \mathcal{N}(\mu_\phi(X_i), \Sigma_\phi(X_i))} [\log \mathcal{N}(f_\theta(Z), \sigma^2 I)] \\ &= \mathbb{E}_{Z \sim \mathcal{N}(\mu_\phi(X_i), \Sigma_\phi(X_i))} \left[-\frac{1}{2} (X_i - f_\theta(Z))^\top (\sigma^2 I)^{-1} (X_i - f_\theta(Z)) \right. \\ &\quad \left. - \frac{1}{2} \log \left((2\pi)^k |\sigma^2 I| \right) \right] \\ &= -\frac{1}{2\sigma^2} \mathbb{E}_{Z \sim \mathcal{N}(\mu_\phi(X_i), \Sigma_\phi(X_i))} [\|X_i - f_\theta(Z)\|^2] - \frac{k}{2} \log(2\pi\sigma^2) \\ &= -\frac{1}{2\sigma^2} \mathbb{E}_{\varepsilon \sim \mathcal{N}(0, I)} \left[\left\| X_i - f_\theta \left(\mu_\phi(X_i) + \sqrt{\Sigma_\phi(X_i)} \varepsilon \right) \right\|^2 \right] - \frac{k}{2} \log(2\pi\sigma^2).\end{aligned}$$

Using the reparametrization trick simplifies the expectation term, and makes able the gradient of this

first term be directly calculated. The second term also can be calculated as the following.

$$D_{\text{KL}}(q_\phi(z | X_i) || r_\lambda(z)) = \frac{1}{2} \left(\text{tr} \left(\text{diag}(\lambda_2)^{-1} \Sigma_\phi(X_i) \right) + (\lambda_1 - \mu_\phi(X_i))^\top \text{diag}(\lambda_2)^{-1} (\lambda_1 - \mu_\phi(X_i)) - k + \log \left(\frac{\det(\text{diag}(\lambda_2))}{\det(\Sigma_\phi(X_i))} \right) \right)$$

The gradient of the second term can also be directly calculated, so we can obtain the gradients via backpropagation.

Problem 4

(a) Let $p_A = (p_{A1}, p_{A2}, p_{A3})$ and $p_B = (p_{B1}, p_{B2}, p_{B3})$. Then

$$\mathbb{E}_{p_A, p_B} [\text{points for } B] = p_{A1}p_{B2} + p_{A2}p_{B3} + p_{A3}p_{B1} - p_{A1}p_{B3} - p_{A2}p_{B1} - p_{A3}p_{B2}.$$

Suppose $p_A^* = (p_{A1}^*, p_{A2}^*, p_{A3}^*)$, $p_B^* = (p_{B1}^*, p_{B2}^*, p_{B3}^*)$ is the solution for the given problem. Then we have

$$\begin{aligned} p_{A1}^*p_{B2} + p_{A2}^*p_{B3} + p_{A3}^*p_{B1} - p_{A1}^*p_{B3} - p_{A2}^*p_{B1} - p_{A3}^*p_{B2} \\ \leq p_{A1}^*p_{B2}^* + p_{A2}^*p_{B3}^* + p_{A3}^*p_{B1}^* - p_{A1}^*p_{B3}^* - p_{A2}^*p_{B1}^* - p_{A3}^*p_{B2}^* \\ \leq p_{A1}p_{B2}^* + p_{A2}p_{B3}^* + p_{A3}p_{B1}^* - p_{A1}p_{B3}^* - p_{A2}p_{B1}^* - p_{A3}p_{B2}^*. \end{aligned}$$

for all $p_A, p_B \in \Delta^3$. If $p_A^* = p_B^* = \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$, all three terms are 0, so it is a solution of the problem.

Now we should show that this is the only solution for the problem. Suppose $p_A^* \neq \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$ and generally let $p_{A1}^* < p_{A2}^*$. Now substitute $p_A = \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$ and $p_B = (0, 0, 1)$, then

$$\begin{aligned} p_{A1}^*p_{B2} + p_{A2}^*p_{B3} + p_{A3}^*p_{B1} - p_{A1}^*p_{B3} - p_{A2}^*p_{B1} - p_{A3}^*p_{B2} &= p_{A2}^* - p_{A1}^* > 0 \\ p_{A1}p_{B2}^* + p_{A2}p_{B3}^* + p_{A3}p_{B1}^* - p_{A1}p_{B3}^* - p_{A2}p_{B1}^* - p_{A3}p_{B2}^* &= 0 \end{aligned}$$

so the inequality becomes false. Similarly, suppose $p_B^* \neq \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$ and generally let $p_{B1}^* < p_{B2}^*$. Now substitute $p_A = (0, 0, 1)$ and $p_B = \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$, then

$$\begin{aligned} p_{A1}^*p_{B2} + p_{A2}^*p_{B3} + p_{A3}^*p_{B1} - p_{A1}^*p_{B3} - p_{A2}^*p_{B1} - p_{A3}^*p_{B2} &= 0 \\ p_{A1}p_{B2}^* + p_{A2}p_{B3}^* + p_{A3}p_{B1}^* - p_{A1}p_{B3}^* - p_{A2}p_{B1}^* - p_{A3}p_{B2}^* &= p_{B1}^* - p_{B2}^* < 0 \end{aligned}$$

so the inequality also becomes false. Therefore always $p_A^* = p_B^* = \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$, so it is the unique solution for the problem.

(b) If B chooses p_B as given, the expected points for B is always 0 regardless of A , so A can choose any

strategy. However, if B chooses strategies other than $p_B = \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right)$, choosing any strategy may not be optimal for A . Choosing $p_B = (1, 0, 0), (0, 1, 0), (0, 0, 1)$ results in $\mathbb{E}_{p_A, p_B}[\text{points for } B] > 0$ each when A chooses strategies such that $p_{A3} > p_{A2}, p_{A1} > p_{A3}, p_{A2} > p_{A1}$.

Problem 3

```
[1]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import transforms

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
batch_size = 128
```

```
[2]: '''
Step 1:
'''

test_val_dataset = datasets.MNIST(root='./mnist_data/',
                                  train=False,
                                  transform=transforms.ToTensor(),
                                  download=True)

test_dataset, validation_dataset = \
    torch.utils.data.random_split(test_val_dataset, [5000, 5000])

# KMNIST dataset, only need test dataset
anomaly_dataset = datasets.KMNIST(root='./kmnist_data/',
                                  train=False,
                                  transform=transforms.ToTensor(),
                                  download=True)
```

```
[3]: '''
Step 2:
'''

# Define prior distribution
class Logistic(torch.distributions.Distribution):
    def __init__(self):
        super(Logistic, self).__init__()

    def log_prob(self, x):
        return -(F.softplus(x) + F.softplus(-x))

    def sample(self, size):
        z = torch.distributions.Uniform(0., 1.).sample(size).to(device)
        return torch.log(z) - torch.log(1. - z)

# Implement coupling layer
class Coupling(nn.Module):
    def __init__(self, in_out_dim, mid_dim, hidden, mask_config):
        super(Coupling, self).__init__()
        self.mask_config = mask_config

        self.in_block = \
            nn.Sequential(nn.Linear(in_out_dim//2, mid_dim), nn.ReLU())
        self.mid_block = \
```

```

        nn.ModuleList([nn.Sequential(nn.Linear(mid_dim, mid_dim), nn.ReLU())
                        for _ in range(hidden - 1)])
    self.out_block = nn.Linear(mid_dim, in_out_dim//2)

def forward(self, x, reverse=False):
    [B, W] = list(x.size())
    x = x.reshape((B, W//2, 2))
    if self.mask_config:
        on, off = x[:, :, 0], x[:, :, 1]
    else:
        off, on = x[:, :, 0], x[:, :, 1]

    off_ = self.in_block(off)
    for i in range(len(self.mid_block)):
        off_ = self.mid_block[i](off_)
    shift = self.out_block(off_)

    if reverse:
        on = on - shift
    else:
        on = on + shift

    if self.mask_config:
        x = torch.stack((on, off), dim=2)
    else:
        x = torch.stack((off, on), dim=2)
    return x.reshape((B, W))

class Scaling(nn.Module):
    def __init__(self, dim):
        super(Scaling, self).__init__()
        self.scale = nn.Parameter(torch.zeros((1, dim)))

    def forward(self, x, reverse=False):
        log_det_J = torch.sum(self.scale)
        if reverse:
            x = x * torch.exp(-self.scale)
        else:
            x = x * torch.exp(self.scale)
        return x, log_det_J

class NICE(nn.Module):
    def __init__(self, in_out_dim, mid_dim, hidden,
                 mask_config=1.0, coupling=4):
        super(NICE, self).__init__()
        self.prior = Logistic()
        self.in_out_dim = in_out_dim

        self.coupling = nn.ModuleList([
            Coupling(in_out_dim=in_out_dim,
                    mid_dim=mid_dim,
                    hidden=hidden,
                    mask_config=(mask_config+i)%2) \

```

```

        for i in range(coupling)])

    self.scaling = Scaling(in_out_dim)

    def g(self, z):
        x, _ = self.scaling(z, reverse=True)
        for i in reversed(range(len(self.coupling))):
            x = self.coupling[i](x, reverse=True)
        return x

    def f(self, x):
        for i in range(len(self.coupling)):
            x = self.coupling[i](x)
        z, log_det_J = self.scaling(x)
        return z, log_det_J

    def log_prob(self, x):
        z, log_det_J = self.f(x)
        log_ll = torch.sum(self.prior.log_prob(z), dim=1)
        return log_ll + log_det_J

    def sample(self, size):
        z = self.prior.sample((size, self.in_out_dim)).to(device)
        return self.g(z)

    def forward(self, x):
        return self.log_prob(x)

```

```

[4]: '''
      Step 3: Load the pretrained model
      '''
      nice = NICE(in_out_dim=784, mid_dim=1000, hidden=5).to(device)
      nice.load_state_dict(torch.load('nice.pt', map_location=device))

```

/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages/torch/distributions/distribution.py:44: UserWarning: <class '_main_.Logistic'> does not define `arg_constraints`. Please set `arg_constraints = {}` or initialize the distribution with `validate_args=False` to turn off validation.

warnings.warn(f'{self.__class__} does not define `arg_constraints`. ' +

```

[4]: <All keys matched successfully>

```

```

[5]: '''
      Step 4: Calculate standard deviation by using validation set
      '''
      validation_loader = torch.utils.data.DataLoader(
          dataset=validation_dataset, batch_size=batch_size)
      likelihood_list = []
      for batch, (images, _) in enumerate(validation_loader):
          likelihood_list += nice(images.view(-1, 784).to(device)).tolist()

      # calculate standard deviation

```

```
import statistics
std = statistics.stdev(likelihood_list)
mean = statistics.mean(likelihood_list)
threshold = mean - 3*std
```

```
[6]: '''
Step 5: Anomaly detection (mnist)
'''
test_loader = torch.utils.data.DataLoader(
    dataset=test_dataset, batch_size=batch_size)
count = 0
for batch, (images, _) in enumerate(test_loader):
    likelihood = nice(images.view(-1, 784).to(device))
    count += torch.sum(likelihood < threshold).item()

print(f'{count} type I errors among {len(test_dataset)} data')
```

103 type I errors among 5000 data

```
[7]: '''
Step 6: Anomaly detection (kmnist)
'''
anomaly_loader = torch.utils.data.DataLoader(
    dataset=anomaly_dataset, batch_size=batch_size)
count = 0
for batch, (images, _) in enumerate(anomaly_loader):
    likelihood = nice(images.view(-1, 784).to(device))
    count += torch.sum(likelihood > threshold).item()

print(f'{count} type II error samong {len(anomaly_dataset)} data')
```

15 type II error samong 10000 data